



# Fuel Deep Dive



Storage




Storage

Contract



Wallet




Receipt

Contract



Predicate




# Three Ethereum Issues

- 25 TPS (142 TPS max)
- State Bloat
- Only 1 source address, only 1 destination address

# Account Model

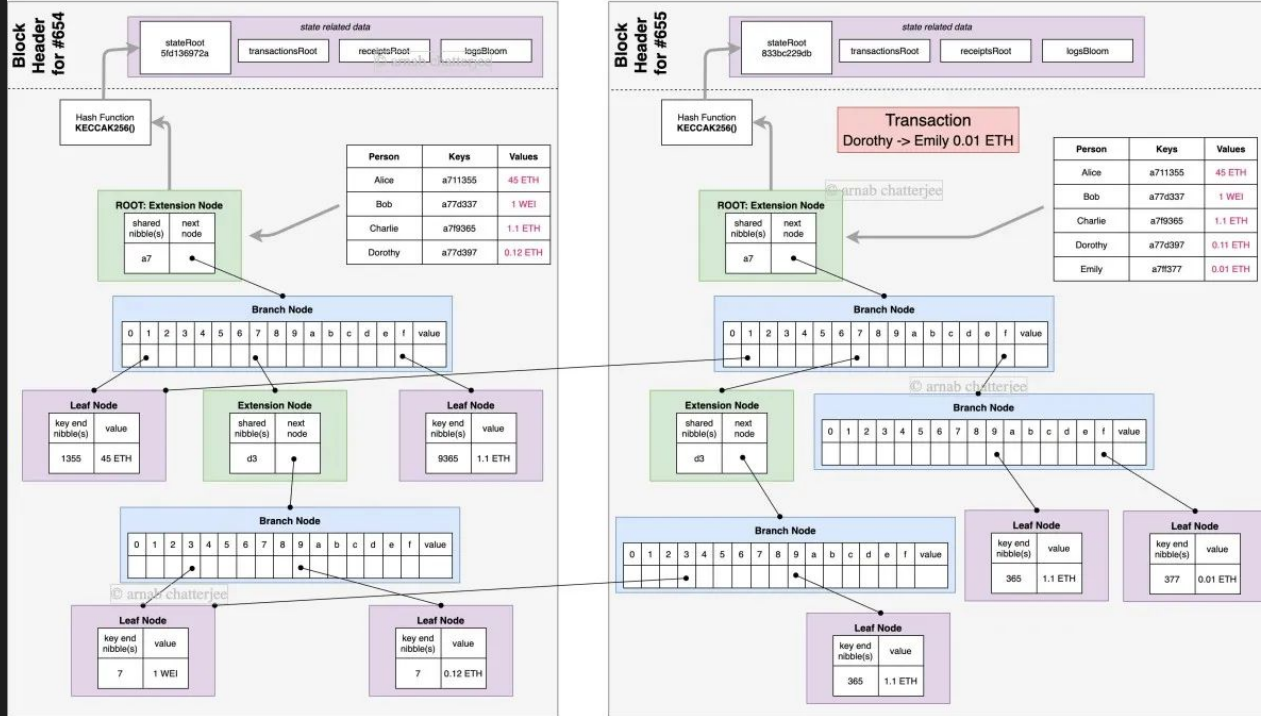
- Ownership is represented by a balance
- Transactions update the associated account balances

*“The world state (state), is a mapping between addresses (160-bit identifiers) and account states (a data structure serialised as RLP, see Appendix B). Though not stored on the blockchain, it is assumed that the implementation will maintain this mapping in a modified Merkle Patricia tree”*



# Account Model

## Merkle Patricia Trie Representation of State Data across Blocks in Ethereum



# Account Model

In a nutshell:

- The state tree is large database of all accounts
- Light nodes don't have to store the whole tree
- They can just query the account balance from a full node and get reasonable assurances with a Merkle proof

# Account Model

But what about parallelism?

CPU 1	CPU 2	CPU 3
TX A		
TX B		
TX C		
TX D		
TX E		

CPU 1	CPU 2	CPU 3
TX A	TX B	TX C
	TX D	TX E



# Account Model

Some transactions need to be processed serially.

TX A  
Sent 10 to Alice

TX B  
Sent the same 10  
to Bob

# Account Model

We can solve this with state access lists

TX A  
Sent 10 from Alice  
to Bob

Access list:  
Alice's balance  
Bob's balance

TX B  
Sent 10 from Alice  
to Uniswap

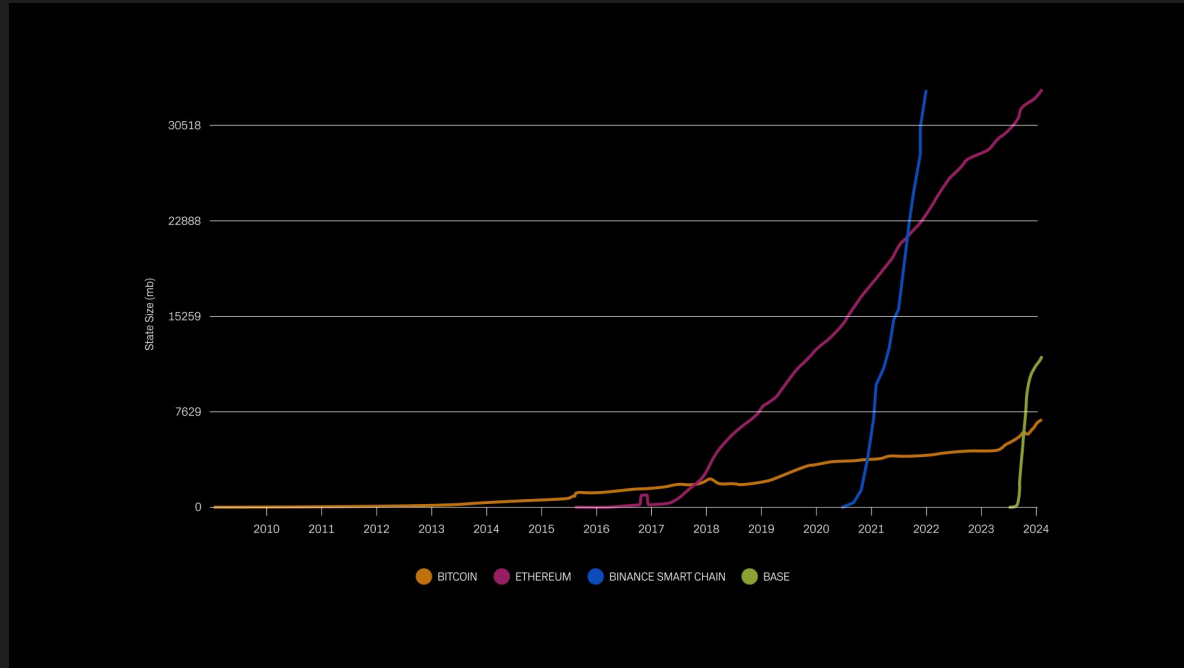
Access list:  
Alice's balance  
Uniswap balance


TX C  
Sent 10 from  
Uniswap to Bob

Access list:  
Bob's balance  
Uniswap balance

# Account Model

We still have a state bloat problem






I THREW OUT THE  
STATE TREE!  
JUST RUN A FULL NODE  
TO KNOW THE STATE!


DON'T I NEED A  
DATA CENTER  
FOR THAT?

# Ethereum state size over time



**Can we do something else than  
Accounts?**








# UTXO Model

## Unspent Transaction Outputs

- Ownership is represented by a set of Unspent Transaction Outputs
- Transactions spend UTXOs and create new ones as output
- No need to store accounts
- Easy to parallelize
- Many inputs, many outputs



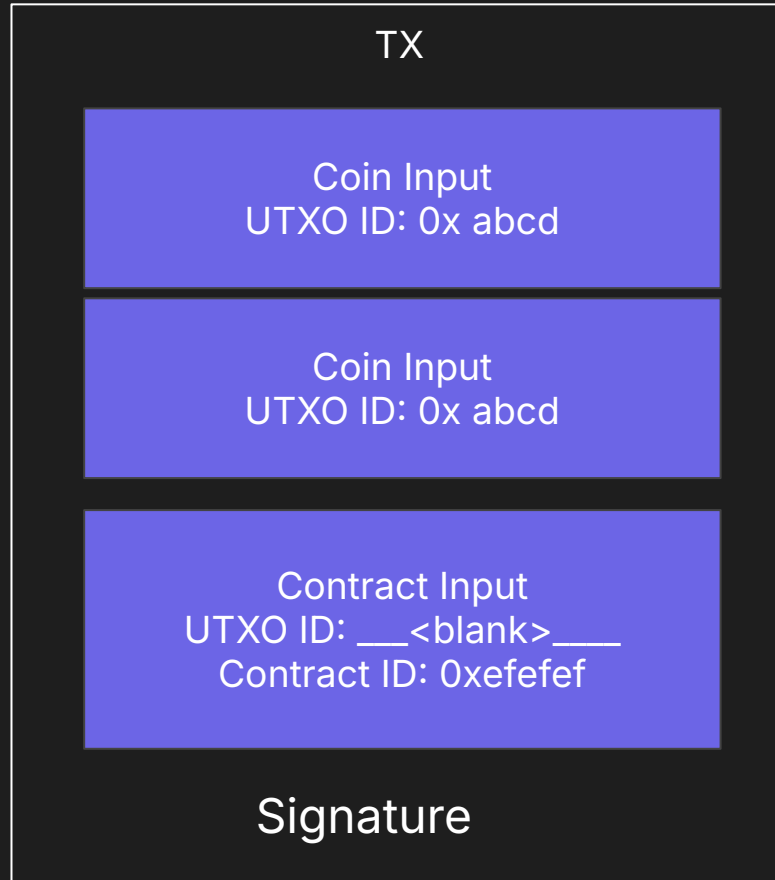
# UTXO Model

- Keeping track of all your UTXOs can be complex
- The Cardano Problem
  - UTXOs are deterministic and can only be spent once
  - But then, how do users of a dapp agree who gets to spend which?



# UTXO Model


I want to use ANY  
Contract UTXO that  
matches this ID



# The Virtual Machine

# The FuelVM

- 64 bit
- Register based
- Byte indexed Stack and Heap model



# The FuelVM

On FuelVM:

- Hashing
- Even more Signature Verification
- All assets are native!

On EVM:

- Hashing
- Signature Verification
- ETH is native, everything else is a contract

# Native Assets

- MINT and BURN instructions
  - `asset_id = sha256((contract_id, sub_id))`
  - SRC-20
- 
- Better DevEx: less boilerplate to make new assets
  - Better Performance: less logic for the VM to process
  - Better Security: no bugs or inconsistencies in transfer

# Executions Contexts

- Scripts
- Predicates
- Contracts



# Scripts

- Entry points for transactions (if you're doing more than just transfers)
- Ephemeral (no state)
- Can call as many contracts as you like

Scripts are essentially "contract launchers"

# Scripts

```
script;
```

```
fn main() {
```

```
    let my_contract = abi(  
        MyContract,  
        MY_CONTRACT_ID  
    );
```

```
    my_contract.my_function();
```

```
}
```

# Scripts

```
script;
```

```
fn main() {
```

```
    let my_contract = abi(  
        MyContract,  
        MY_CONTRACT_ID  
    );
```

```
    my_contract.first_function();
```

```
    my_contract.second_function();
```

```
}
```

# Scripts

Replaces  
Uniswap-style router  
contract

```
script;
```

```
fn main() {
```

```
    let amm = abi(MyAmm, MY_CONTRACT_ID);  
    amm.read_from_pools();  
    do_some_calculations();  
    amm.swap(A, B);  
    amm.swap(B, C);  
    ensure_final_output_is_sufficient();
```

```
}
```

# Predicates

- Can own coins
- not stored on chain
- stateless spending condition

# Predicates

Coin held by a normal account



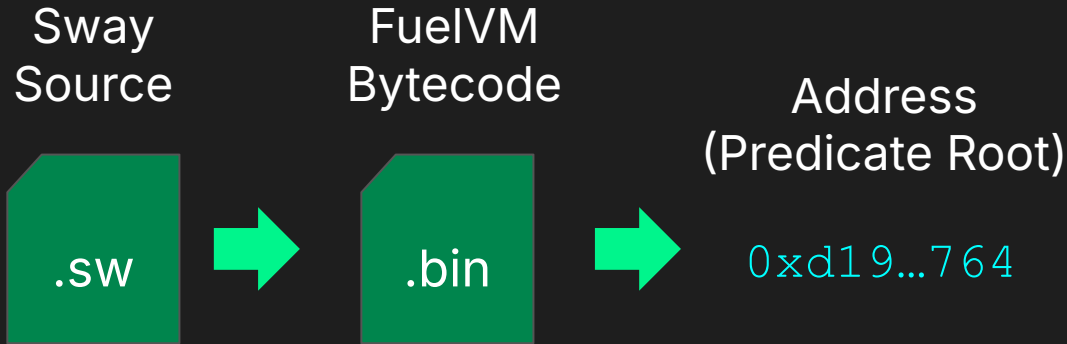
Can be spent with a signature

Coin held by a predicate root



Can be spent by executing a  
predicate matching the root and  
returning true

# Predicates



All funds sent to a predicate root can be spent if one can execute the predicate and return true

# Predicates

```
predicate;
```

```
fn main() -> bool {  
    Time::now() > UNLOCK_TIMESTAMP  
}
```

**Anybody** can spend the funds after this date



# Predicates

```
predicate;
```

```
fn main(sig: b256) -> bool {  
    validate_signature(sig)  
}
```

# Contracts

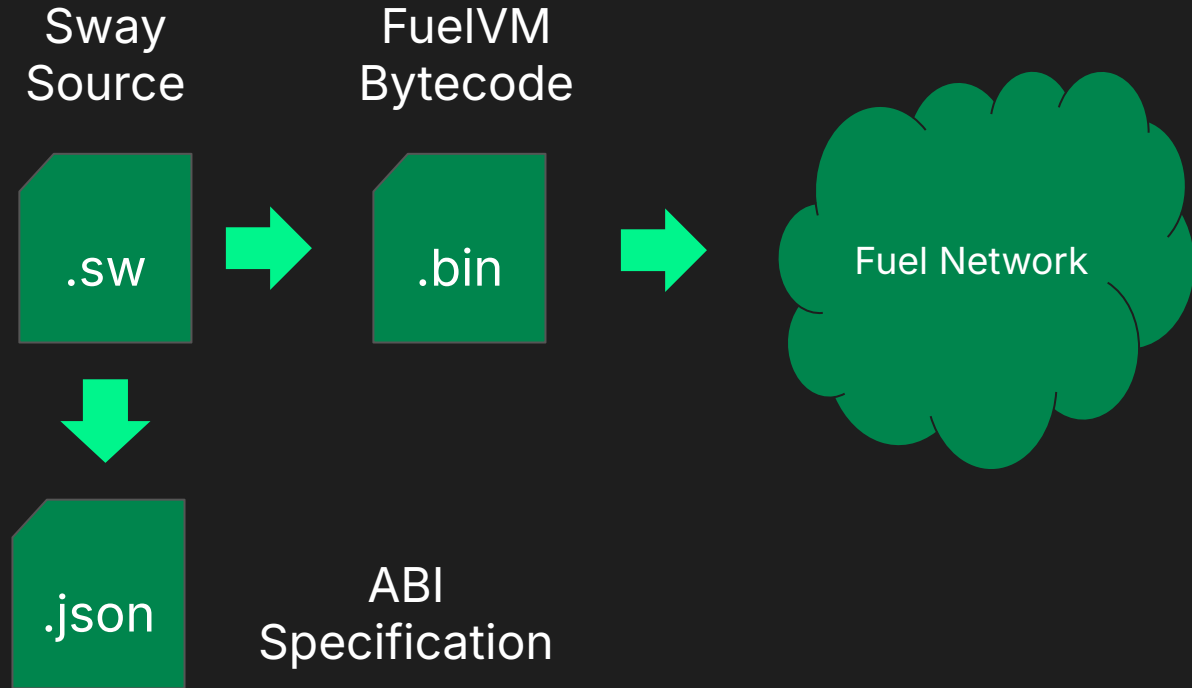
Smart contracts with their own:

- Coins (contracts can own assets)
- ABI
- Execution context
- Storage context

# Contract ABI

```
contract;  
  
impl Contract {  
    fn foobar() {  
        log("hello");  
    }  
}
```

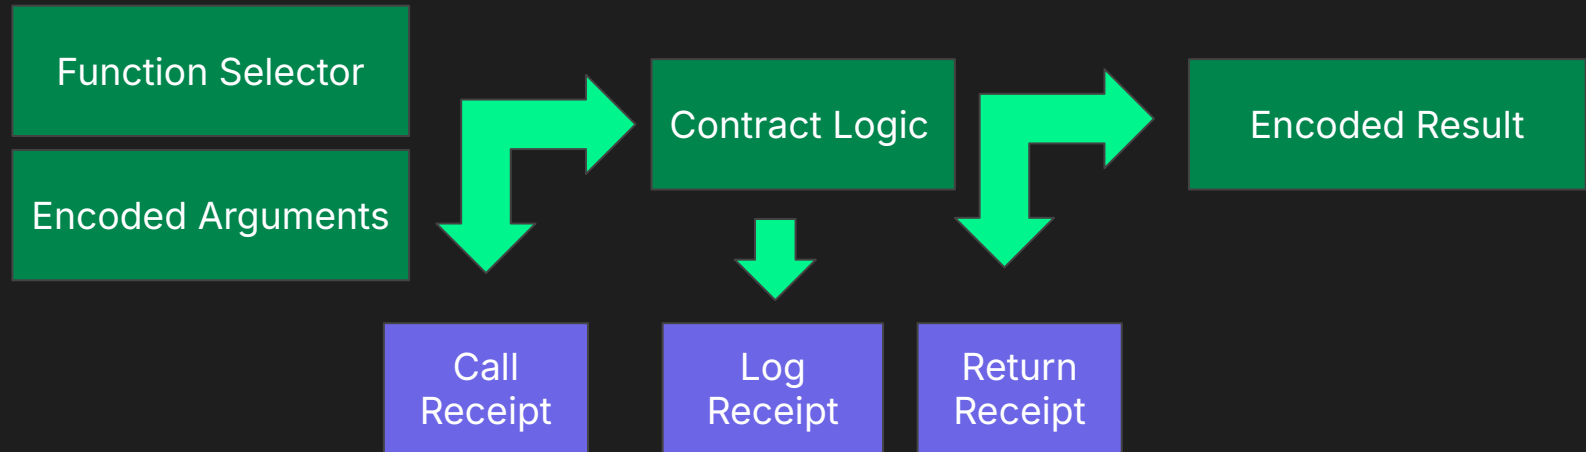
# Contract ABI



# Contract ABI

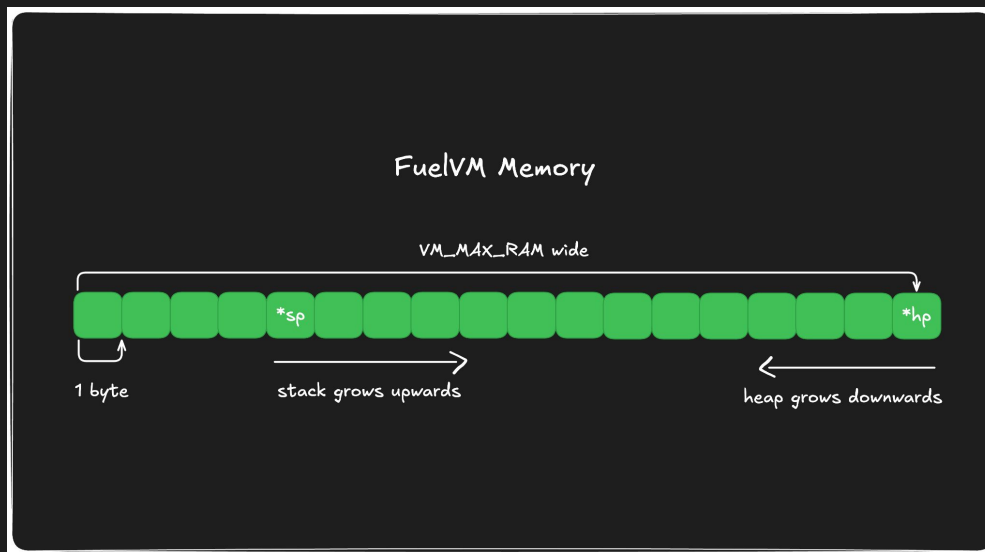
```
...  
"functions": [  
  {  
    "name": "foobar",  
    "inputs": [],  
    "output": "2e3...f5d",  
    "attributes": null  
  },  
],  
...
```

# Contract ABI




# Memory

- `ssp`: bottom of the current writable stack area
- `sp`: on top of the stack
- `hp`: below the current bottom of the heap



# Memory


## Memory Write Policies



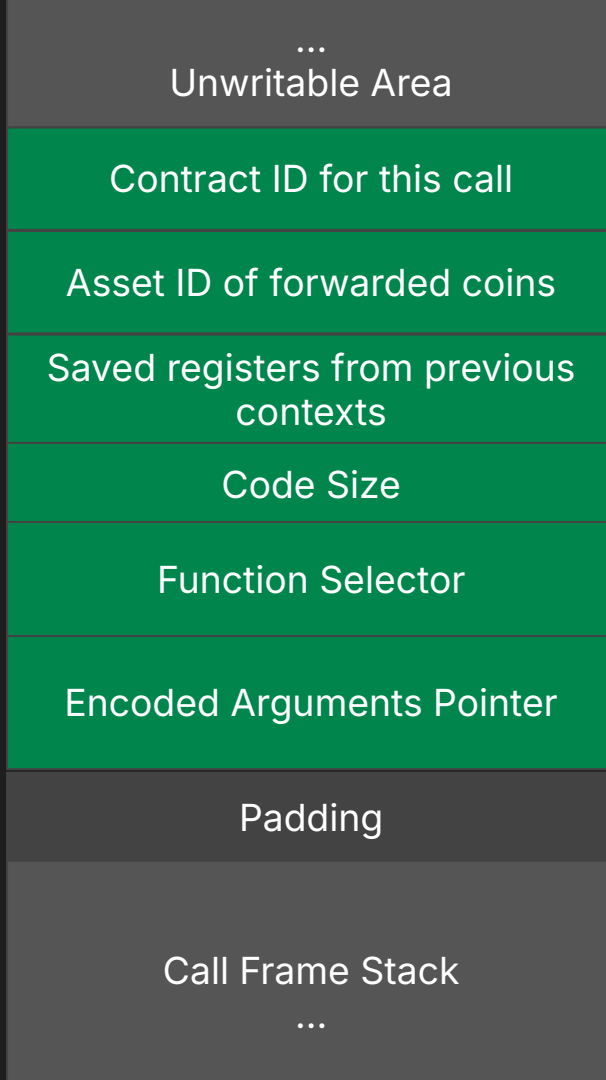


# Memory

## Memory Read Policies



# Memory



A typical call frame

# Storage

- Persistent storage
- Every contract has its own independent storage context
- $2^{256}$  slots of 256 bits (sparse)


```
storage {  
    foo: StorageMap<u64, u64> = StorageMap{},  
}
```

1	2	5
42	43	46


sha256(  
 (1, "storage\_0")  
)

sha256(  
 (2,  
 "storage\_0")  
)


sha256(  
 (5, "storage\_0")  
)



**Putting it all together**



Storage



Storage

Contract



Wallet



Receipt

Contract



Predicate